


# Real-Time Streaming Showdown



Flink

**Sherif Behna**  
Data Architect  
Hikari Data inc.

# About me




# What is Stream Processing ?

- Stream : Continuous sequence of events over time
- Stream processing : Processing these events to get insights

# Typical Use Cases

- Real-time analytics
- Event-driven applications
- Sensor data processing (IoT)
- Business monitoring and alerting
- Anomaly and fraud detection

# Typical Architecture



# Do you really need "real-time"

- Real-time means milliseconds to a few seconds (for example : 5s end-to-end)
- Leads to extra complexity (vs batch or "slow real-time")
- Requires better monitoring and alerting
- Is there a business case for it ?

# Demo

### Stock Ticker Generator

Single event    Multiple events

Stock Symbol  
AAPL

Price (\$)  
e.g., 150.50

Send Stock Data

#### Recent Events

Symbol	Price	Timestamp
AAPL	\$125.00	2/21/2026, 4:19:43 PM
AAPL	\$115.00	2/21/2026, 4:19:27 PM
AAPL	\$109.00	2/21/2026, 4:19:11 PM
AAPL	\$101.00	2/21/2026, 4:19:09 PM
AAPL	\$100.00	2/21/2026, 4:19:04 PM

## Stock Alerts Dashboard

Real-time alerts from Spark & Flink processors

FLINK ALERTS: 3    FLINK WINDOWED: 3    SPARK ALERTS: 4    SPARK WINDOWED: 4


### Flink Processor

TIME	SYMBOL	PRICE	VARIATION %
4:19:43 PM	AAPL	\$125.00	+8.7%
4:19:27 PM	AAPL	\$115.00	+5.5%
4:19:11 PM	AAPL	\$109.00	+7.92%


### Flink Windowed Processor

TIME	SYMBOL	PRICE	VARIATION %
4:19:44 PM	AAPL	\$115.00	+5.5%
4:19:27 PM	AAPL	\$109.00	+8.46%
4:19:27 PM	AAPL	\$100.50	N/A

# State management under the hood



# More advanced state management




# Apache Spark



- **Micro-batch processing** - Divides stream into small batches
- **APIs:**
  - DStream API (legacy)
  - Structured Streaming (DataFrame API)
  - Spark SQL integration
- **Additional libraries:** MLlib, GraphX, Spark Connect, Pandas Spark
- **Mature ecosystem**

# Spark Architecture



# Spark Code Snippet



```
Dataset<Row> lines = spark
    .readStream()
    .format("socket")
    .option("host", SOCKET_HOST)
    .option("port", SOCKET_PORT)
    .load();

Dataset<StockData> stockData = lines
    .flatMap(new JsonParser(), Encoders.bean(StockData.class));

Dataset<String> alerts = stockData
    .groupByKey(
        (MapFunction<StockData, String>) StockData::getSymbol,
        Encoders.STRING()
    )
    .mapGroupsWithState(
        new PriceChangeDetector(),
        Encoders.bean(StockState.class),
        Encoders.STRING(),
        GroupStateTimeout.NoTimeout()
    )
    .filter((FilterFunction<String>) alert -> alert != null && !alert.isEmpty());


StreamingQuery query = alerts
    .writeStream()
    .outputMode("update")
    .format("console")
    .start();
```

# Apache Flink



- **Designed for real-time stream processing**
- **Processes events independently (true streaming)**
- **APIs:**
  - DataStream API (low-level, event-by-event processing)
  - Table API (relational operations)
  - Flink SQL
- **Additional libraries:** Flink CDC, Flink ML, Flink CEP, Flink Agents
- **Mature ecosystem (less than Spark but mature enough)**

# Flink Architecture



# Flink Code Sample




```
final StreamExecutionEnvironment env = StreamExecutionEnvironment.getExecutionEnvironment();

DataStream<String> socketStream = env.socketTextStream(SOCKET_HOST, SOCKET_PORT);


socketStream
    .flatMap(new JsonParser())
    .keyBy(StockData::getSymbol)
    .flatMap(new PriceChangeDetector())
    .print();

env.execute("Flink Stock Price Alert Processor");
```


# Micro-batch vs Continuous



**Flink**



# Performance



Source: <https://yahooenq.tumblr.com/post/135321837876/benchmarking-streaming-computation-engines-at>

# Expected Features of a Streaming Platform

Feature	Apache Spark	Flink
Supports multiple sources and sinks	✓	✓
Multiple connectors	✓	✓
Integration with monitoring and management tools	✓	✓
Language support (streaming)	Java, Scala, Python, SQL, R	Java, Scala, Python
Scalability (distributed)	✓	✓

# Stateless Operators

Feature	Apache Spark	Flink
Transformations (map / flatMap)	✓	✓
Filters	✓	✓
Projections (select)	✓	✓
Join with static data	✓	✓
Partitioning	✓	✓

# Stateful Operators

Feature	 Apache Spark	 Flink
Grouping	✓	✓
Aggregations (sum, count, avg, custom)	✓	✓
Windows (tumbling and sliding)	✓	✓
Stream to stream join	✓	✓
Sessionalizing	✓	✓

# Handling the Pitfalls of Streaming

Feature	Apache Spark	Flink
Time management (event time vs processing time)	✓	✓
Late-arriving data (Watermarks)	✓	✓
Fine-grained state management	Limited	✓
Delivery semantics (at-most once, at-least once, exactly once)	at-least once	exactly once
Fault tolerance	limited	✓
Handle backpressure	limited	✓

# Bread and Butter Considerations

Feature	 Apache Spark	 Flink
Learning curve	Moderate	Steep
Product maturity	High	Medium
Ecosystem maturity	High	Medium
Vendor support	High	Low
Scalability	Medium	High

# And the winner is...

 Apache Spark	 Flink
Mature and stable ecosystem	Low latency and predictable performance
More vendor / tooling support	Responsive event-driven apps
Data engineering / ML use cases	Robust state management and fault tolerance

# Q & A

Repo : <https://github.com/shbehna/streaming-spark-flink>

